Fynd AI Intern Take Home Assessment Report

Task 1: Rating Prediction via Prompting

- I loaded a reviews dataset and used it to test how well a large language model can predict star ratings.

- I designed three different prompt styles: zero-shot, few-shot, and self-correction based prompt.
- The model was asked to read review text and assign a rating from 1 to 5 stars.
- I ran all prompts on the same test data to keep comparison fair.
- I measured performance using accuracy, precision, recall, and F1-score.
- I also checked how many valid predictions the model was able to give.
- Based on the evaluation results, I compared all prompt strategies.
- The best performing prompt was selected using macro F1-score.

Task 2: Two-Dashboard AI Feedback System

- I built a web application with two dashboards that are deployed.
- Both dashboards read and write to the same CSV file called submissions.csv.
- The user dashboard is public and allows users to select a rating and write a review.
- When the user submits the form, an AI response is generated and shown to the user.
- The user input is saved into the CSV file for storage.
- The admin dashboard is internal and shows all submissions in a live list.
- It displays the user rating and the review text.
- It also shows AI generated summary of the review and recommended actions.
- I also added simple analytics like total number of submissions.